

Distributed Denial of Service (DDoS): Attacks and Defenses

I. Pustogarov



Forwarding-Loop Attacks in Content Delivery Networks

NDSS 2016

<http://www.icir.org/vern/papers/cdn-loops.NDSS16.pdf>

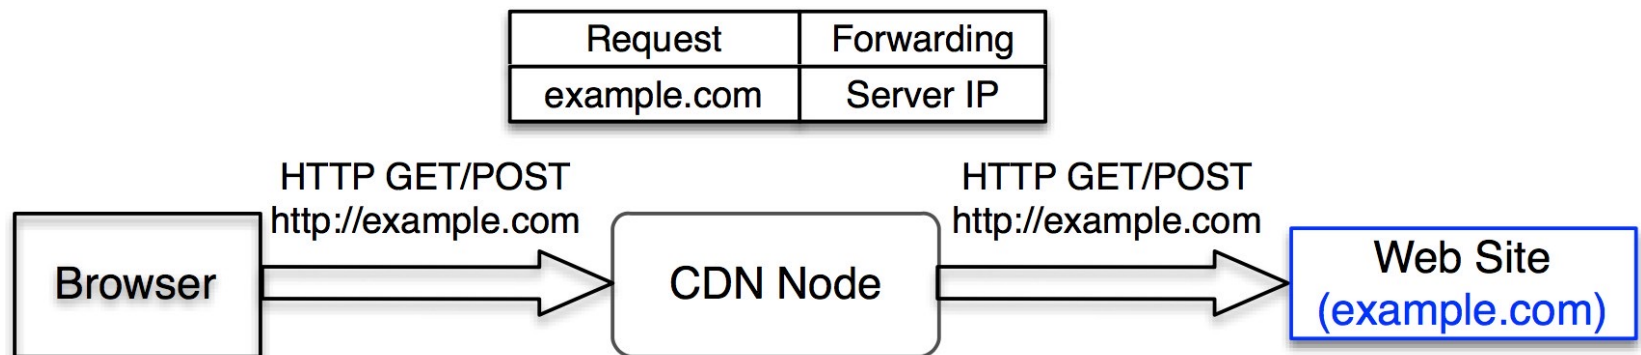
Summary

- Present “forwarding loop” attacks that threaten CDN availability
- A case that highlights the danger of allowing cross-organization, user-controlled (untrusted) policies without centralized administration
- Measured 16 popular CDNs and find all of them are vulnerable to such attacks
- Vendors have acknowledged the problem

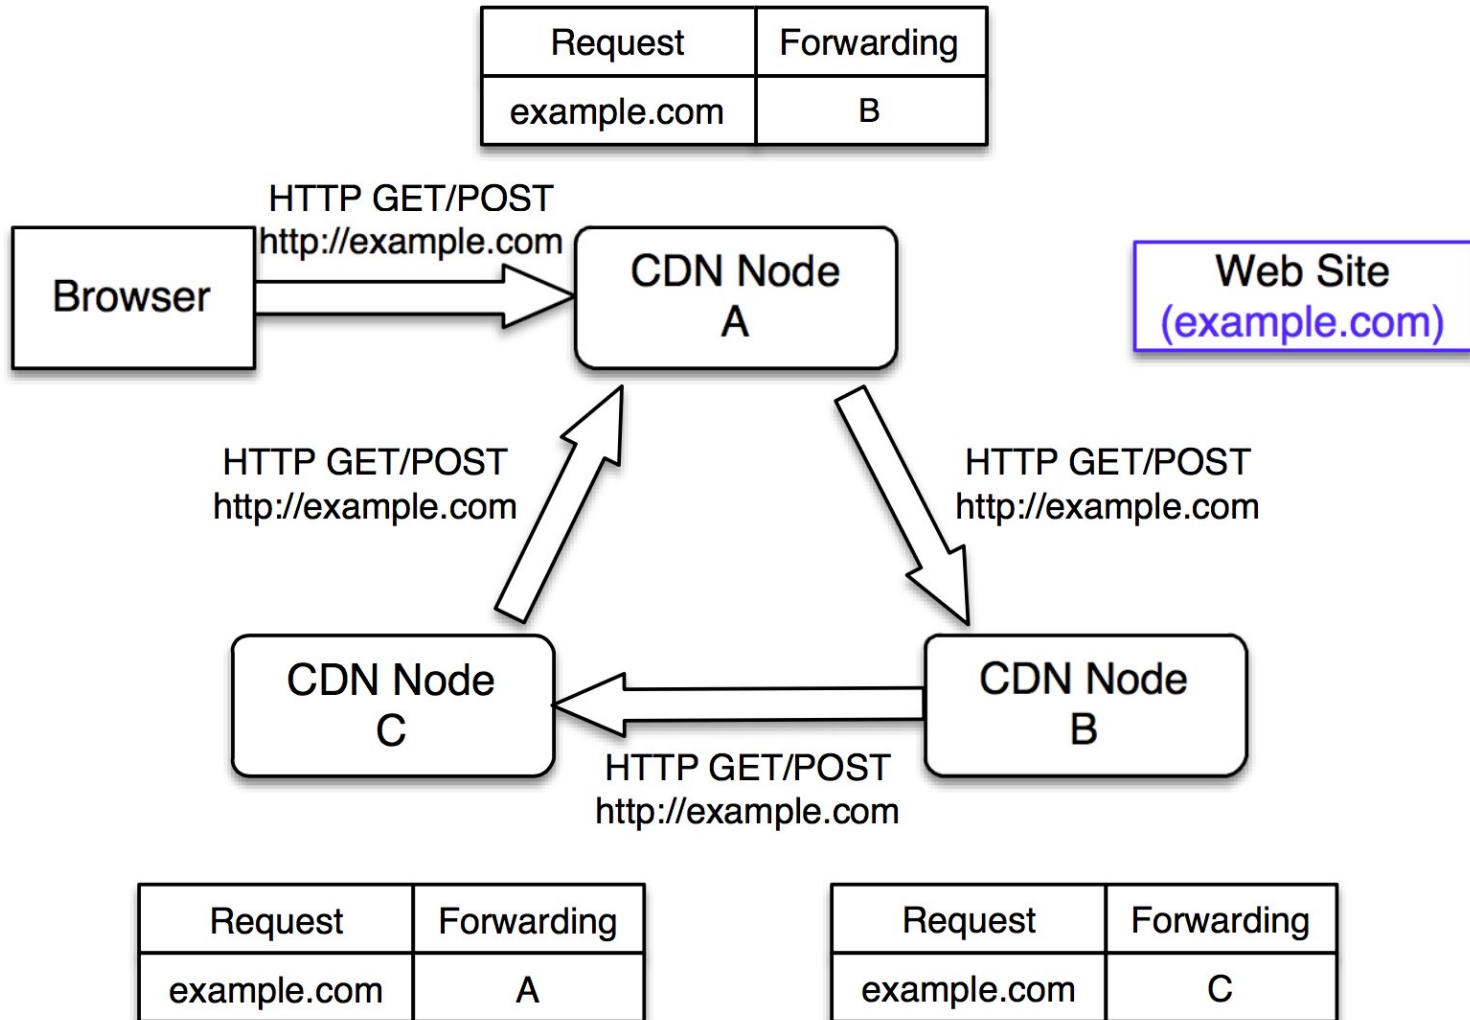
Tested CDNs



Normal CDN forwarding rule



CDN Forward loop creation



Attack practicality

- Cost : All 16 CDNs provide free or free-trial account
- Can remain anonymous: 11/16 CDNs only require an email address
- Some CDNs agreed this attack is severe

Defenses

- Unifying and standardizing a loop-detection header (“Via” as recommended by RFC)
- Interim defenses
 - Obfuscating self-defined loop-detection headers
 - Monitoring and rate-limiting
 - Constraint on forwarding destination



Mirai botnet and DDoS attacks

Understanding the Mirai Botnet

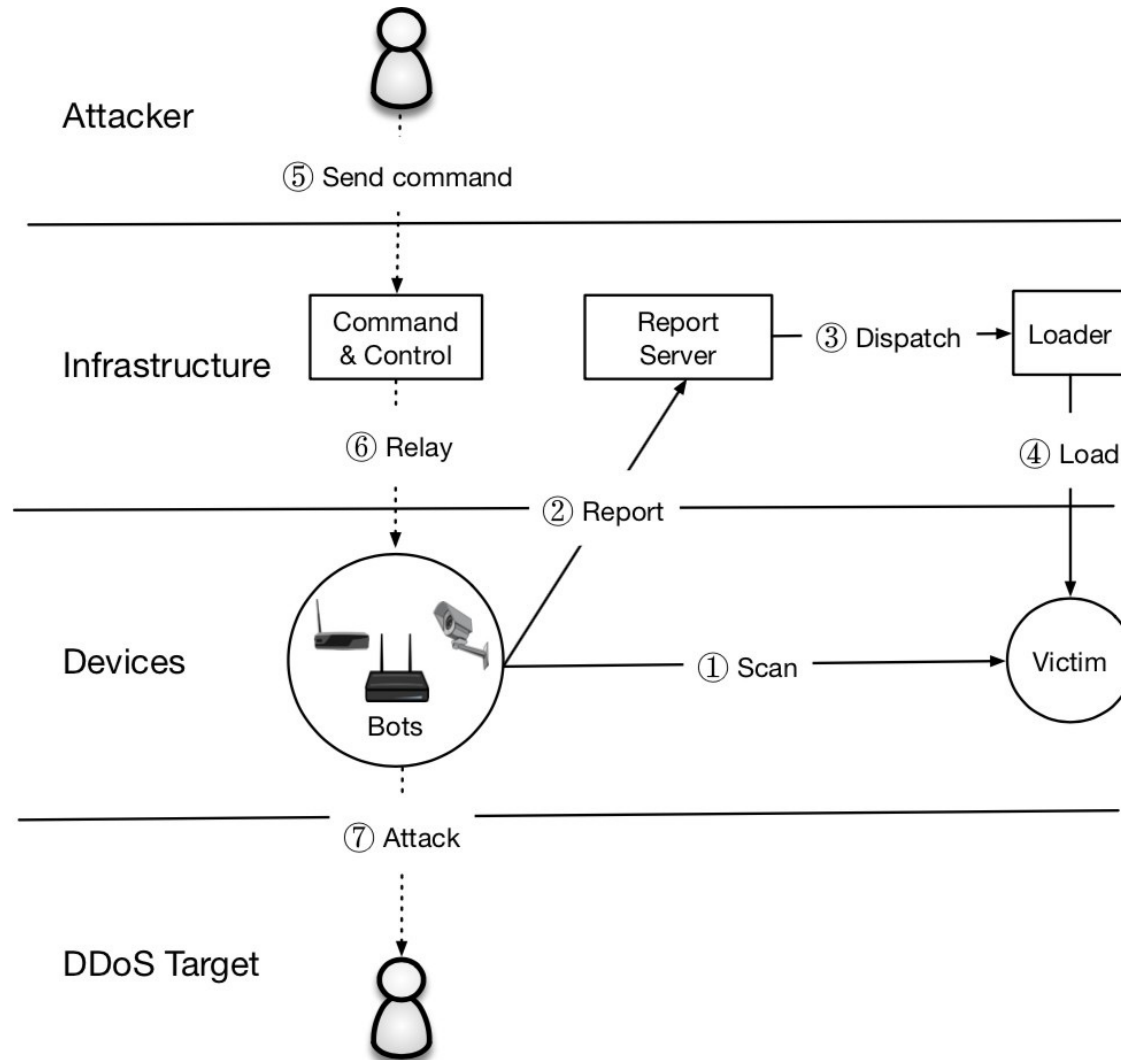
(USENIX Security 2017)

<https://www.usenix.org/system/files/conference/usenixsecurity17/sec17-antonakakis.pdf>

- Starting in September 2016, several massive DDoS attacks were launched from the Mirai botnet
- Victims: Krebs on Security (620 Gbps), OVH (1Tbps), and Dyn (affected: GitHub, Twitter, Reddit, Netflix, Airbnb)
- Abused hundreds of thousands of some of the Internet's least powerful hosts — IoT devices

- The Mirai botnet, composed primarily of embedded and IoT devices: DVRs, IP cameras, routers, and printers
- Infected nearly 65,000 IoT devices in its first 20 hours before reaching a steady state population of 200,000–300,000 infections; **peak infection: 600,000**
- Features: efficient spreading based on Internet-wide scanning, rampant use **of insecure default passwords** in IoT products, and the insight that keeping the botnet's behavior simple would allow it to infect many **heterogeneous devices**

Mirai Operation



Mirai Operation

- Starts with: *the rapid scanning* phase “statelessly” sent TCP SYN probes to pseudorandom IPv4 addresses, **excluding those in a hard-coded IP blacklist**, on Telnet TCP ports 23 and 2323
- If a potential victim is found, starts a *brute-force login* phase: a Telnet connection using 10 username and password pairs selected randomly from a pre-configured list of 62 credentials.
- At the first successful login, Mirai sent the victim IP and associated credentials to a hard-coded *report server*.
- A separate *loader program* asynchronously infected these vulnerable devices by logging in, determining the underlying system environment, and finally, downloading and executing architecture-specific malware.

Mirai Operation

- After a successful infection, Mirai attempted to **conceal its presence** by deleting the downloaded binary and obfuscating its process name in a pseudorandom alphanumeric string.
- As a consequence, Mirai infections **did not persist** across system **reboots**.
- It also killed other processes bound to TCP/22 or TCP/23, as well as processes associated with competing infections, including other Mirai variants

DDoS attack types

Attack Type	Attacks	Targets	Class
HTTP flood	2,736	1,035	A
UDP-PLAIN flood	2,542	1,278	V
UDP flood	2,440	1,479	V
ACK flood	2,173	875	S
SYN flood	1,935	764	S
GRE-IP flood	994	587	A
ACK-STOMP flood	830	359	S
VSE flood	809	550	A
DNS flood	417	173	A
GRE-ETH flood	318	210	A

Table 9: **C2 Attack Commands**—Mirai launched 15,194 attacks between September 27, 2016–February 28, 2017. These include [A]pplication-layer attacks, [V]olumetric attacks, and TCP [S]tate exhaustion, all of which are equally prevalent.

DDoS attack types

- Only 2.8% of Mirai attack commands relied on bandwidth amplification
- Amplification attacks make up 74% of attacks issued by DDoS-for-hire booter services

Timeline

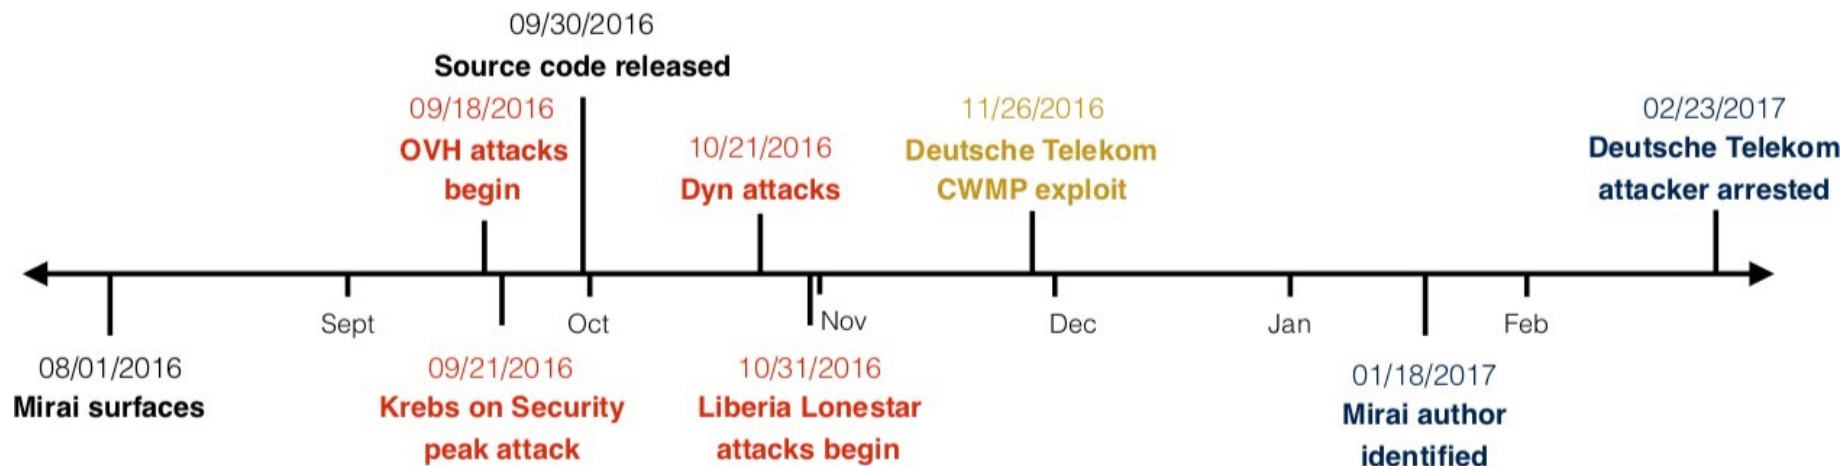


Figure 1: **Mirai Timeline**—Major attacks (red), exploits (yellow), and events (black) related to the Mirai botnet.



“Botz-4-Sale: Surviving Organized
Attacks That Mimic Flash Crowds”, Kandula et al.

Kill-bots – a solution at the target

- Motivations
 - Bots can easily mimic flash crowds (differ in intent – but in content)
 - Requests originate from diverse geographic locations
 - Can authentication, puzzles help?
- Kill-bots
 - Kernel extension to protect web servers
 - Combine two functions:
 - Authentication
 - Admission control
 - Relies on: automated Turing tests (ATTs)
 - E.g.: CAPTCHAs
 - Also, relies on: attackers must launch a “fast” DDoS to succeed
 - Why so?



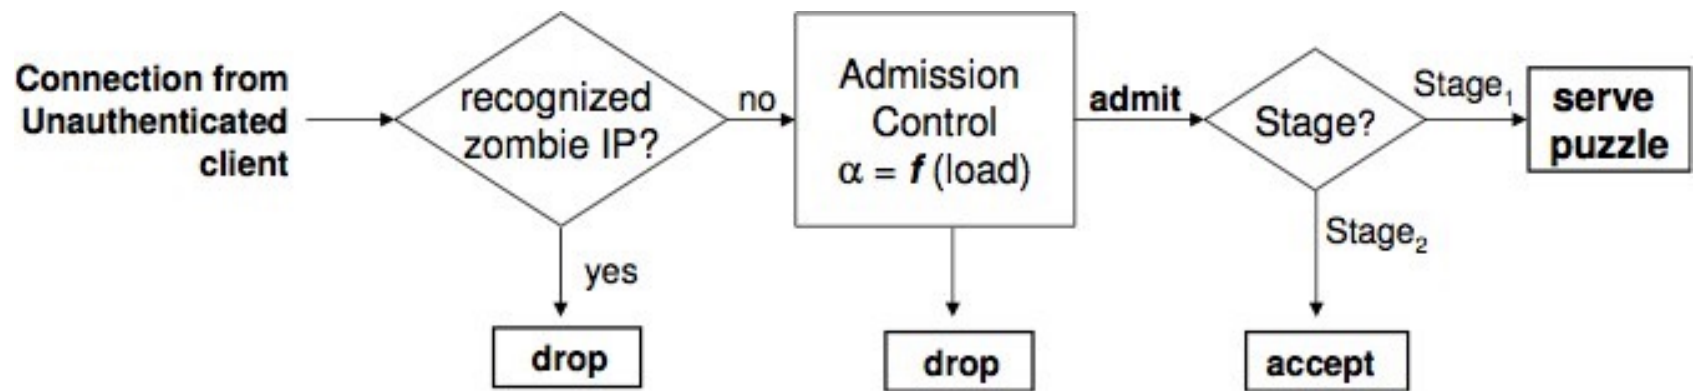
Kill-bots: Authentication

- Activated when the server is overloaded
- Two stages
- Stage 1:
 - Each new session must solve an ATT
 - Human users may solve ATTs – if they want service
 - Bots will be “unable” to solve ATTs
 - SYN-cookies are used to address spoofed IPs
 - Bloom filter is used to count how often an IP address failed to solve a test
 - The server discards requests from a client if the number of its unsolved tests exceeds a given threshold (e.g., 32)

Kill-bots: Authentication Stage 2

- Kill-bots switches to Stage2 after the set of detected zombie IP addresses stabilizes
 - The filter does not learn any new bad IP addresses
- ATTs are no longer served
- Kill-bots relies solely on the Bloom filter to drop requests from malicious clients (from stage 1)
- This allows legitimate users who cannot, or do not want to solve ATTs access to the server despite the ongoing attack

Kill-bots: an overview



Activating Kill-bot authentication

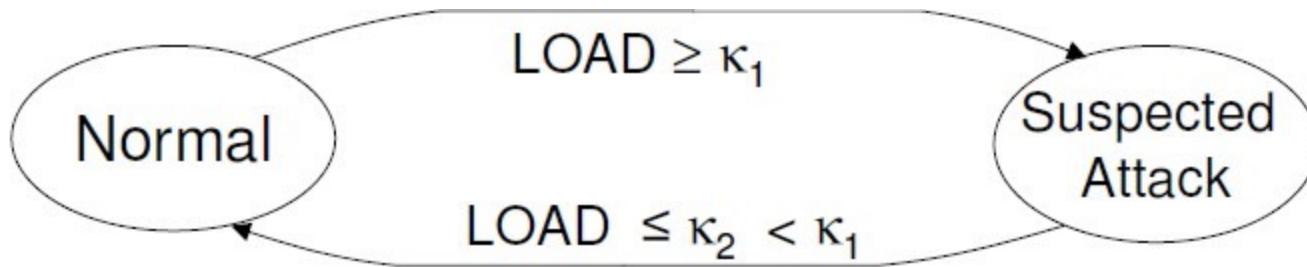


Figure 2: A Kill-Bots server transitions between NORMAL and SUSPECTED_ATTACK modes based on server load.

Security analysis of Kill-bots

1. Social engineering attacks
 - Use unsuspecting users to solve ATTs – effective against email / account creation ATTs
 - But in Kill-bots: ATT answers expire (after 4 minutes)
 - i.e.: continuous authentication
2. Bloom-filter pollution: fails due to SYN cookies (no IP spoofing)
3. Copy attacks: limit in-progress requests to a low threshold (e.g., 8)
4. Database with answers to all ATTs
 - Add new ATTs to frustrate the attacker
5. Breaking ATTs
 - Design better ones